

## Dynamic Programming Algorithms

Uses the Bellman equations to define iterative algorithms for both policy evaluation and control

- Policy evaluation defines state value function  $v_\pi$  for a particular policy  $\pi$
- Policy evaluation improves a particular policy by modifying it to make a better policy. We can continuously improve the policy until it is no longer possible to improve, which means the last policy must be equal to the optimal policy

### 1. Initialization

$V(s) \in \mathbb{R}$  and  $\pi(s) \in \mathcal{A}$  arbitrarily for all  $s \in \mathcal{S}$

### 2. Iterative Policy Evaluation

Input  $\pi$ , the policy to be evaluated

$V \leftarrow \vec{0}, V' \leftarrow \vec{0}$

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$V'(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |V'(s) - V(s)|)$

$V \leftarrow V'$

Until  $\Delta \leftarrow \theta$

Output  $V \approx v_\pi$

### 3. Policy Improvement

*policy - stable*  $\leftarrow$  true

For each  $s \in \mathcal{S}$ :

*old - action*  $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$

If *old - action*  $\neq \pi(s)$ , then *policy - stable* = false

If *policy - stable*, then stop and return  $V \approx v_*$  and  $\pi \approx \pi_*$ ; else go to 2

## Generalized Policy Iteration

### 1. Value iteration

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of the estimation.

Initialize  $V(s)$ , for all  $s \in S^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

$V \leftarrow V'$

Until  $\Delta \leftarrow \theta$

Output a deterministic policy  $\pi \approx \pi_*$ , such that:

$\pi(s) = \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$

## 2. Alternatives

- Monte Carlo method: averaging the value on a set of results from the policy  $\pi$
- Bootstrapping: using the previous value estimate to improve the current value estimate
- Brute-force search: evaluating every possible deterministic policy one at the time and selecting the one with the highest value