

Learning parameterized policies

In reinforcement learning, learning parameterized policies means directly optimizing a policy $\pi(a|s; \theta)$ with respect to its parameters θ , instead of indirectly via value functions.

This is the foundation of policy gradient methods, which are especially powerful when combined with function approximation (e.g., neural networks).

A **parameterized policy** is a function $\pi(a|s; \theta)$ that maps a state to a probability distribution over actions, with parameters θ (e.g., weights of a neural network).

Examples:

- **Discrete actions:** Softmax over linear preferences
- **Continuous actions:** Gaussian policy with $\mu(s; \theta), \sigma(s; \theta)$

Policy Gradient Theorem

The **policy gradient** gives the direction to adjust θ to improve expected return:

$$\nabla_{\theta} J(\theta) = E_{\pi}[\nabla_{\theta} \log \pi(A_t|S_t; \theta) \cdot q_{\pi}(S_t, A_t)]$$

We use the gradient to perform **stochastic gradient ascent**:

$$\theta \leftarrow \theta + \alpha \cdot \nabla_{\theta} J(\theta)$$

Practical Algorithm: REINFORCE

A Monte Carlo method using full return G_t :

$$\theta \leftarrow \theta + \alpha \cdot G_t \cdot \nabla_{\theta} \log \pi(A_t|S_t; \theta)$$

Can be high variance, so usually paired with:

- **Baselines** (e.g., subtract $v_{\pi}(s)$) to reduce variance
- **Actor-Critic methods** (learn both policy and value function)

Actor-Critic Methods

Use a **value function critic** $\hat{v}(s; w)$ or **advantage function** $\hat{A}(s, a)$ to guide the policy update:

$$\theta \leftarrow \theta + \alpha \cdot \hat{A}(s, a) \cdot \nabla_{\theta} \log \pi(a|s; \theta)$$

- **Actor:** the policy $\pi(a|s; \theta)$
- **Critic:** the value estimator (can also be neural network)

Policy Gradient for Continuing Tasks

In **continuing tasks** (no terminal states), we aim to **optimize performance over an infinite horizon** — without relying on discounting or episode boundaries. This setting is often more realistic in domains like robotics, process control, or system maintenance.

$$\rho(\theta) = \lim_{t \rightarrow \infty} \mathbb{E}_{\theta}[R_t]$$

Policy Gradient Theorem (Average Reward Form)

The gradient of the average reward with respect to policy parameters is:

$$\nabla_{\theta} \rho(\theta) \propto \sum_s d^{\pi}(s) \sum_a \nabla_{\theta} \pi(a|s; \theta) \cdot q_{\pi}(s, a)$$

Where:

- $d^{\pi}(s)$: stationary distribution under π
- $q_{\pi}(s, a)$: expected total reward **above average** starting from (s, a)

This leads to the familiar stochastic gradient estimate:

$$\nabla_{\theta} \rho(\theta) \approx \mathbb{E}_{\pi} [\nabla_{\theta} \log \pi(A_t | S_t; \theta) \cdot (R_{t+1} + \hat{v}(S_{t+1}) - \hat{v}(S_t))]$$

Implementation in Actor-Critic (Continuing)

You use a **differential TD error** to update both actor and critic:

TD Error (Differential Form):

$$\delta_t = R_t + 1 - \bar{R} + \hat{v}(S_{t+1}) - \hat{v}(S_t)$$

Critic Update:

$$w \leftarrow w + \beta \cdot \delta_t \cdot \nabla_w \hat{v}(S_t; w)$$

Actor Update:

$$\theta \leftarrow \theta + \alpha \cdot \delta_t \cdot \nabla_{\theta} \log \pi(A_t | S_t; \theta)$$

Where:

- \bar{R} : moving estimate of the average reward

Policy Parameterizations

A **parameterized policy** is a function:

$$\pi(a | s; \theta)$$

which gives the probability of selecting action a in state s , controlled by parameters θ .
These parameters can be:

- **Linear weights**
- **Neural network weights**
- **Coefficients in softmax or Gaussian distributions**

Common Forms of Policy Parameterization

1. Softmax (Discrete Actions)

Used when action space is **discrete**:

$$\pi(a | s; \theta) = \frac{e^{h(s,a;\theta)}}{\sum_b e^{h(s,b;\theta)}}$$

- $h(s, a; \theta)$: preference score (can be linear or nonlinear)
- Often used in REINFORCE or Actor-Critic

2. Gaussian (Continuous Actions)

Used when action space is **continuous**:

$$\pi(a | s; \theta) = \mathcal{N}(\mu(s; \theta), \sigma^2(s; \theta))$$

- Mean and std dev are outputs of a neural network
- Used in robotics, control, etc.

Gradient of Log Policy

Key property used in **policy gradient methods**:

$$\nabla_{\theta} \log \pi(a \mid s; \theta)$$

This term allows you to compute how the probability of choosing an action changes as parameters are adjusted. It's used in the update:

$$\theta \leftarrow \theta + \alpha \cdot G_t \cdot \nabla_{\theta} \log \pi(A_t \mid S_t; \theta)$$

or with TD error δ_t in Actor-Critic.